BigHammer.ai



# Make failure cheap and loud.



4 guardrails that cut wasted **Databricks spend** on failed jobs. All configuration. No migration.

- Timeouts
- Retries
- Pre-checks
- Cost alerts

BigHammer.ai

# 1 Timeouts set from real runtime

Set it from each job's own history, for example **2x the median.** A 45-minute job then stops at 90, not at sunrise.






## No timeout means no stop.

Default per the Databricks Jobs API docs.

BigHammer.ai

# 2 Retries with a ceiling and an interval

Retries help with flaky sources. With bad code they just **repeat the bill.** Cap them, space them out, and agree when a failure needs a human.



```
"max_retries": 2,
```

↖ Default is 0. Set -1 and it retries forever.

```
"min_retry_interval_millis": 600000
```

↖ Leave it out and retries fire at once.

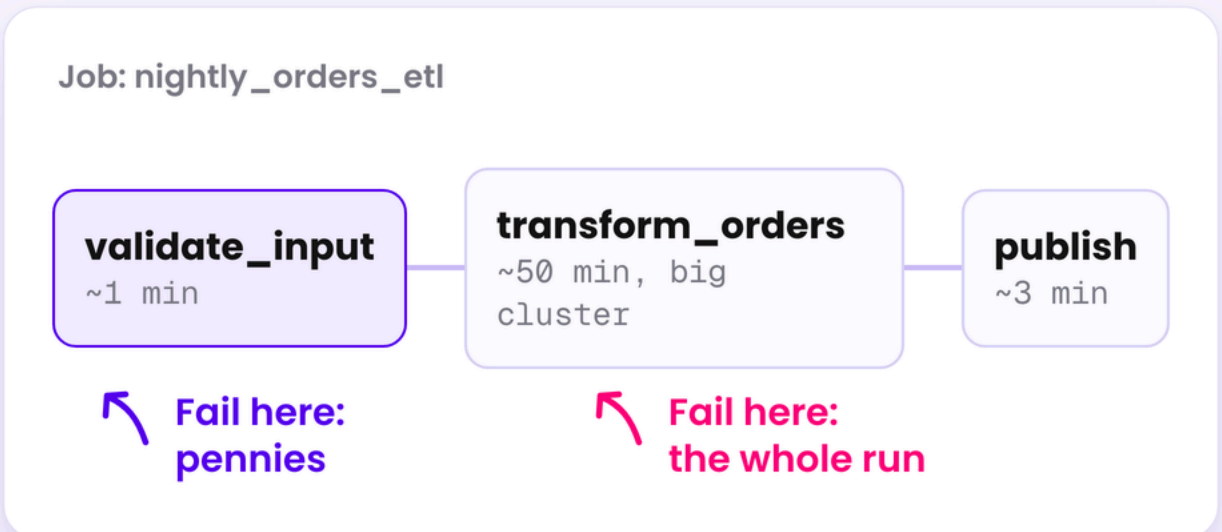
## Two retries, ten minutes apart. Then a human.

Defaults per the Databricks Jobs API docs.

# 3 A cheap check before the expensive work



Is the input there? Is it the right shape? Is it non-empty? **Failing at minute one costs pennies.** Failing at minute fifty costs the whole run.




**Put the gate at the front.**

BigHammer.ai

# 4 Alerts that include the cost and the owner



Put the failed-run cost and a named owner in the alert. Failure stops being invisible to the people who **pay for it.**

`nightly_orders_etl` failed — Acknowledged

`nightly_orders_etl` failed after 2 retries and burned **$312** of compute.
Owner: Data Platform team — Fixed

**“Failed” gets acknowledged. “Burned $312” gets fixed.**

BigHammer.ai

# Failed runs are one of the quickest Databricks savings to find. See the rest live.

## Cut your Databricks bill by up to 75%

Free live masterclass
Thu 15 Oct, 12:00 PM ET
5 PM UK
With **Srinath Reddy**, founder and CEO,
and **Richard Lawrence**, BigHammer.ai



**webinar.bighammerai.com**

Free to join. Bring your failed-job list.