

# AI AGENTS CAN MOVE YOUR DATA. CAN THEY TRANSFORM IT?



Terry Dhariwal



# WHAT THE BENCHMARKS ACTUALLY SAY

Everyone has a demo where an AI agent builds a pipeline in minutes.

The public benchmarks tell a **more useful story**: which part of the work is solved, and which part isn't yet.

# AI ON DATA WORK, BY THE NUMBERS

**96%**

## **EXTRACT AND LOAD**

moving data from source to target

**66%**

## **DBT PROJECT TASKS**

best score, Spider 2.0-DBT

**33%**

## **FULL TRANSFORMATION**

raw data to correct models, end to end

Two different benchmarks. Spider 2.0-DBT scores narrower dbt project tasks; ELT-Bench-Verified asks the agent to build the full transformation layer end to end.

1

# MOVING DATA IS NEARLY SOLVED



In the ELT-Bench paper, extract and load success rose from **37% to 96%** when the agent moved to a newer model.

One setup, one benchmark, but a clear signal: use agents here now.

2

# MEANING IS THE HARD PART

Transformation is where business rules live: what counts as revenue, which customer is active, how two systems reconcile.

Agents get it right about **one time in three**. That's progress. It isn't autopilot.



3

# SO SPLIT THE WORK

Only **33%** of developers say they trust AI accuracy.

Stack Overflow Developer Survey 2025

**AI** for the repetitive, checkable steps

**Fixed rules** wherever results must be predictable

**People** own the business logic and review what ships



# USE AGENTS WHERE THEY'RE STRONG.

Keep human judgement where they're not. That's how you get the speed without shipping wrong numbers to the board.

Which part of your pipeline would you trust an agent with today?