

AI AGENTS CAN MOVE YOUR DATA. CAN THEY TRANSFORM IT?



Terry Dhariwal

WHAT THE BENCHMARKS ACTUALLY SAY

Everyone has a demo where an AI agent builds a pipeline in minutes.

The public benchmarks tell a **more useful story**: which part of the work is solved, and which part isn't yet.

AI ON DATA WORK, BY THE NUMBERS

96%

EXTRACT AND LOAD

moving data from source to target

66%

DBT PROJECT TASKS

best score, Spider 2.0-DBT

33%

FULL TRANSFORMATION

raw data to correct models, end to end

Two different benchmarks. Spider 2.0-DBT scores narrower dbt project tasks; ELT-Bench-Verified asks the agent to build the full transformation layer end to end.

ELT-Bench-Verified, Mar 2026: one agent setup (Claude Sonnet 4.5 + SWE-agent), 96% / 32.51%. Spider 2.0-DBT leaderboard, best 65.6%.

1

MOVING DATA IS NEARLY SOLVED



In the ELT-Bench paper, extract and load success rose from **37% to 96%** when the agent moved to a newer model.

One setup, one benchmark, but a clear signal: use agents here now.

2

MEANING IS THE HARD PART

Transformation is where business rules live: what counts as revenue, which customer is active, how two systems reconcile.

Agents get it right about **one time in three**. That's progress. It isn't autopilot.



3

SO SPLIT THE WORK

Only **33%** of developers say they trust AI accuracy.

Stack Overflow Developer Survey 2025

AI for the repetitive, checkable steps

Fixed rules wherever results must be predictable

People own the business logic and review what ships



USE AGENTS WHERE THEY'RE STRONG.

Keep human judgement where they're not. That's how you get the speed without shipping wrong numbers to the board.

Which part of your pipeline would you trust an agent with today?